

→ λ_c increases →

Pareto trade-off: NDCG@20 vs CS@20 under varying inference-time penalty λ_c

- TRIER (no content, $\gamma_o=0$)
- TRIER-L (no content, $\gamma_o=0.01$)
- TRIER-C (+content, $\gamma_o=0$)
- PACER-Full (+content, $\gamma_o=0.01$)

